

Benchmark	What it isolates	Best alternative	Ours
RedCode-Exec	<i>granted tool, ungranted target</i>	0% (RBAC / capability token)	99.9%
AgentDojo (live)	<i>prompt-injection ASR</i>	5.6% built-in; Progent 11–17%	$\approx$ 0%
AgentDojo (replay)	<i>injection / arg swap on traces</i>	77.9% below binding	100%
IPI / ATB / AdvBench / MCP / ToolEmu	<i>unauthorized recipient / poisoned tool</i>	0–40% below binding	100%
Business-process (live)	<i>authorized sequence, composite rule</i>	$\approx$ 100% viol. (Progent / CaMeL)	0% viol.
Burst + flow + BFCL detector	<i>authorized bursts & chunked leaks</i>	0% (per-call ladder)	100% / 0% FA
SLEIGHT (twins)	<i>covert intent, same tools/paths</i>	$\approx$ 6% (<i>n</i> -gram)	$\approx$ 98%
Adaptive red-team	<i>exfil / persist / destroy, mutating</i>	0% (RBAC / capability token)	100%
Long-horizon	<i>needle in session's own tool</i>	0–83% (lower rungs)	100% / 0 FA
Cross-session + drift	<i>split spend; baseline poison</i>	session ledger fails; rolling 0.5%	holds / 82%
Mandate quality	<i>incomplete grant (held-out FB)</i>	47.9% FB on τ^2	0.05% FB
iVisor + latency	<i>syscall escape / egress; cost</i>	floor misses BPL; CaMeL minutes	8/8; $\sim$20 μs
Utility cost	<i>clean-task cost (deployable)</i>	ARGUS $\sim$ 5; CaMeL $\sim$ 7; Progent $\sim$ 19	3 pts (grok)
ULB / IBM AML	<i>peer deviation + fan-out</i>	0.86 AUC	0.914; 15/15